

# DEEP LEARNING FOR ARTIFICIAL INTELLIGENCE

3rd Master Course UPC ETSETB TelecomBCN Barcelona. Autumn 2019.



## Instructors



Xavier  
Giró-i-Nieto



Marta R.  
Costa-jussà



Noé  
Casas



Verónica  
Vilaplana



Ramon  
Morros



Javier  
Ruiz



Albert  
Pumarola

## Organizers



UNIVERSITAT POLITÈCNICA  
DE CATALUNYA  
BARCELONATECH



## Supporters



Google Cloud

GitHub Education

+ info: <http://bit.ly/dlai2019>

[\[course site\]](#)



#DLUPC

Day 5 Lecture 2

## Methodology



Javier Ruiz Hidalgo

[javier.ruiz@upc.edu](mailto:javier.ruiz@upc.edu)

Associate Professor

Universitat Politècnica de Catalunya  
Technical University of Catalonia



# Outline

- **Capacity of the network**
  - Underfitting
  - Overfitting
- **Prevent overfitting**
  - Dropout, regularization
- **Data**
  - training, validation, test partitions
  - Augmentation
- **Strategy**

# Outline



# It's all about the data...





Figure extracted from Kevin Zakka's Blog, [Nuts and Bolts of Applying Deep Learning](#), 2016.

# well, not only data...

- Computing power: GPUs



Source: NVIDIA 2017


# well, not only data...

- Computing power: GPUs
- New learning architectures
  - CNN, RNN, LSTM, DBN, GNN, GAN, etc.


<http://www.asimovinstitute.org/neural-network-zoo/>

# End-to-end learning: speech recognition

## Traditional model




## End-to-end learning




This works well given enough labeled (audio, transcript) data.

# End-to-end learning: autonomous driving

## Traditional model



## End-to-end learning




Given the safety-critical requirement of autonomous driving and thus the need for extremely high levels of accuracy, a pure end-to-end approach is still challenging to get to work. End-to-end works only when you have enough (x,y) data to learn function of needed level of complexity.




# Network capacity

- Space of representable functions that a network can potentially learn:
  - Number of layers / parameters



# Generalization

The network needs to **generalize** beyond the training data to work on new data that it has not seen yet






# Underfitting vs Overfitting

- **Overfitting:** network fits training data too well
  - Excessively complicated model
- **Underfitting:** network does not fit training data well enough
  - Excessively simple model

**Both underfitting and overfitting lead to poor predictions on new data and they do not generalize well**


# Underfitting vs Overfitting



# Data partition

How do we measure the generalization instead of how well the network does with the memorized data?

Split your data into two sets: training and test



# Underfitting vs Overfitting




Figure extracted from [Deep Learning](#) by Ian Goodfellow and Yoshua Bengio and Aaron Courville, MIT Press, 2016

# Data partition revisited

- Test set **should not** be used to tune your network
  - Network architecture
  - Number of layers
  - Hyper-parameters
- Failing to do so will overfit the network to your test set!
  - <https://www.kaggle.com/c/higgs-boson/leaderboard>

# Data partition revisited (2)

- Add a **validation** set!



- Lock away your test set and use it only as a last validation step / measure of performance



# The bigger the better?

- Large networks
  - More capacity / More data
  - Prone to overfit
- Smaller networks
  - Lower capacity / Less data
  - Prone to underfit




# The bigger the better?

- The probability of finding a “bad” (high value) local minimum is non-zero for small networks and decreases quickly with network size.
- In large networks, most local minima are equivalent and yield similar performance.
- Struggling to find the global minimum on the training set (as opposed to one of the many good local ones) is not useful in practice and may lead to overfitting.


**Better large capacity networks and prevent overfitting**

# Prevent overfitting

- Early stopping
- Regularization
- Dropout
- Data augmentation
- Parameter sharing
- Adversarial training




# Early stopping



# Weight regularization (1)

- Control the capacity of the network to **prevent overfitting**
- Large weights tend to cause sharp transitions in node functions → large changes in output for small changes in inputs
  - **Penalize the weights** of the nodes in the network
  - **Discourages** learning a more **complex** or flexible **model**



# Weight regularization (2)

- L2-regularization (weight decay):

$$\mathcal{L}_{new} = \mathcal{L} + \frac{\lambda}{2} W^2$$

regularization  
hyper-parameter  
↙


- L1-regularization:

$$\mathcal{L}_{new} = \mathcal{L} + \frac{\lambda}{2} |W|$$

# Weight regularization (3)

- Limit the values of parameters in the network
  - L2 vs L1 regularization

$$\mathcal{L}_{new} = \mathcal{L} + \frac{\lambda}{2} W^2 \quad \mathcal{L}_{new} = \mathcal{L} + \frac{\lambda}{2} |W|$$




# Weight regularization (4)

- L2 regularization heavily penalizes peaky weights and prefers diffuse / low value weights
- L1 regularization leads weights to become sparse (i.e. very close to exactly zero)



# Activation regularization (1)

- Large activations (output values) can also denote overfitting or instability
- Instead of controlling the weights (weight regularization) → control directly the output of nodes (activation regularization)
  - **Penalize the activations** of the nodes in the network
  - **Encourages** networks with low activations




# Activation regularization (2)

- Two common methods for calculating the magnitude of the activation are:
  - Sum of absolute activation values  $\rightarrow$  L1 norm
  - Sum of squared activation values  $\rightarrow$  L2 norm
- L1 norm encourages sparsity
- L2 norm encourages small activations values in general
  
- A hyperparameter is needed to control the amount of penalty in the activation.



# Dropout (1)

- At each training iteration, randomly remove some nodes in the network along with all of their incoming and outgoing connections ([N. Srivastava, 2014](#))



# Dropout (2)

- Why dropout works?
  - Nodes become more insensitive to the weights of the other nodes  $\rightarrow$  more robust.
  - Averaging multiple models  $\rightarrow$  **ensemble**.
  - Training a collection of  $2^n$  thinned networks with parameters sharing




Ensemble of subnetworks


# Dropout (3)

- Every forward pass, network slightly different.
- Reduce co-adaptation between neurons
- More robust features
- Dropout is **removed** in validation/testing

✗ **More iterations** for convergence




(a) Standard Neural Net



(b) After applying dropout.

# Dropout (4)

- Dense-sparse-dense training ([S. Han 2016](#))
  - a. Initial regular training
  - b. Drop connections where weights are under a particular threshold.
  - c. Retrain sparse network to learn weights of important connections.
  - d. Make network dense again and retrain using small learning rate, a step which adds back capacity.



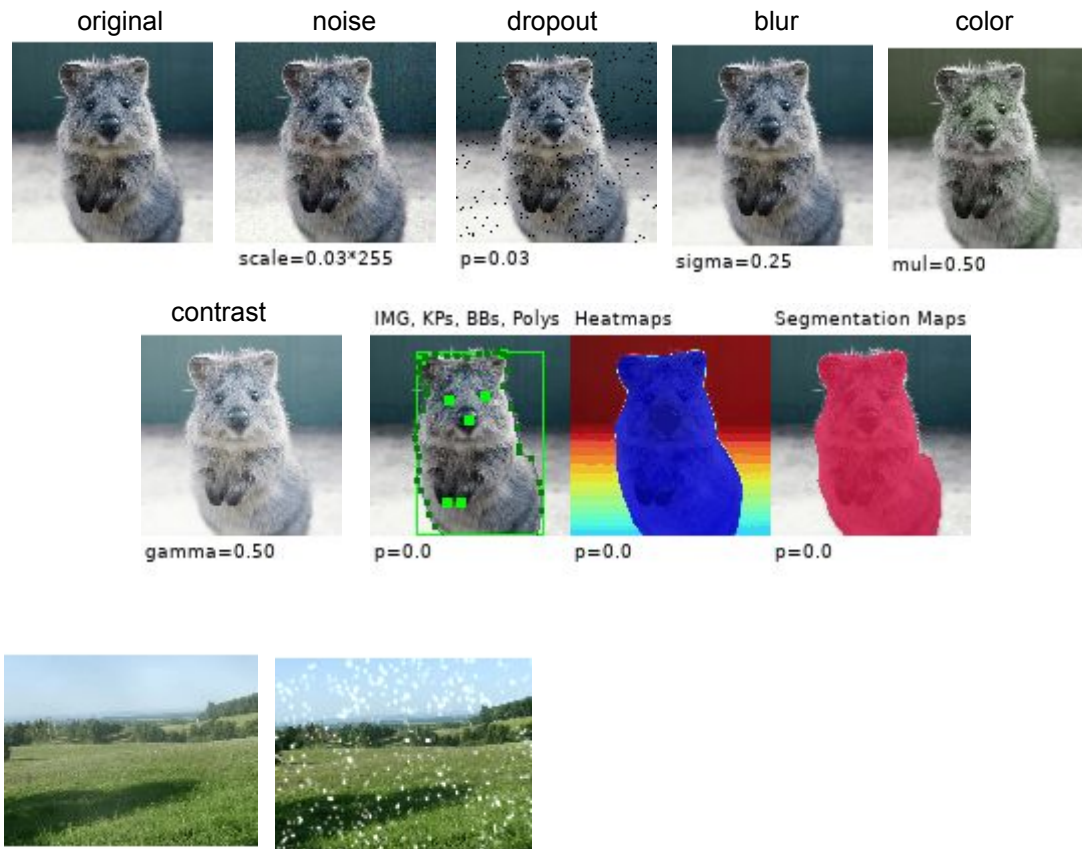
# Data augmentation (1)

- Modify input samples artificially to increase the data size
- On-the-fly while training
  - Inject Noise
  - Transformations
- Not used in testing/validation



# Data augmentation (2): Image



- Noise injection
- Dropout
- Blurs
- Color changes
- Contrast
- Transformations
  - GT transformed!
- Crops, shifts
- Application specific
  - Clouds, snow, etc.





# Data augmentation (3): Audio

- Noise injection
- Shifting time
- Changing pitch
- Changing speed
- Crops
- Loudness
- Mask



# Data augmentation (4): Text

- OCR changes
  - $o \rightarrow 0$
  - $i \rightarrow 1$
- Keyboard changes
  - Substitute character close on keyboards
- Random changes
- Substitute entire words
  - spelling mistakes
  - word2vec distance



# Data augmentation (5)


- Synthetic data: Generate new input samples



A. Palazzi, [Learning to Map Vehicles into Bird's Eye View](#), ICIAP 2017  
DeepGTAV plugin: <https://github.com/ai-tor/DeepGTAV>  
[CARLA Simulator](#).

# Data augmentation (6)

- GANs (Generative Adversarial Networks)



# Parameter sharing

## Multi-task Learning

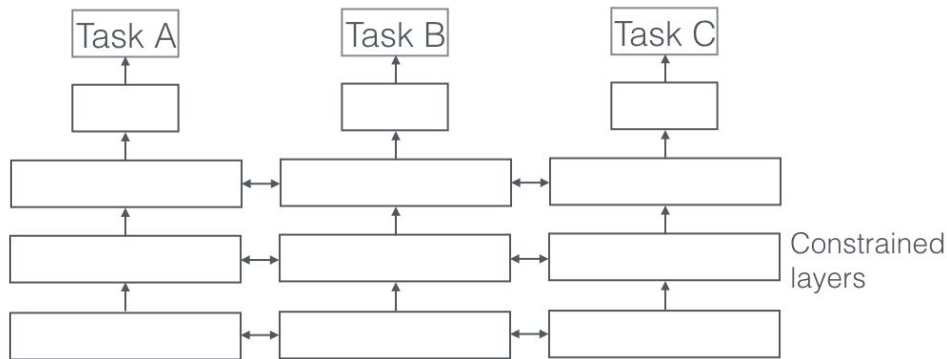




Figure extracted from Sebastian Ruder, [An Overview of Multi-Task Learning in Deep Neural Networks](#), 2017

### FULLY CONNECTED NEURAL NET



### LOCALLY CONNECTED NEURAL NET



CNNs

# Adversarial training

- Search for **adversarial examples** that network misclassifies
  - Human observer cannot tell the difference
  - However, the network can make highly different predictions.



$x$   
 $y = \text{"panda"}$   
w/ 57.7%  
confidence

+ .007 ×



$\text{sign}(\nabla_x J(\theta, x, y))$   
"nematode"  
w/ 8.2%  
confidence


=



$x + \epsilon \text{sign}(\nabla_x J(\theta, x, y))$   
"gibbon"  
w/ 99.3%  
confidence

# Strategy for machine learning (1)

Human-level performance can serve as a very reliable proxy which can be leveraged to determine your next move when training your model.



# Strategy for machine learning (2)

TRAINING 60%	VALIDATION 20%	TEST 20%
-----------------	-------------------	-------------

Human level error . . 1%

Training error . . . 9%

Validation error . . 10%

Test error . . . . 11%



Underfitting



# Strategy for machine learning (3)

TRAINING 60%	VALIDATION 20%	TEST 20%
-----------------	-------------------	-------------

Human level error . . 1%

Training error . . . 1.1%

Validation error . . 10%

Test error . . . . 11%



Overfitting training

# Strategy for machine learning (4)

TRAINING 60%	VALIDATION 20%	TEST 20%
-----------------	-------------------	-------------

Human level error . . 1%

Training error . . . 1.1%

Validation error . . 2%

Test error . . . . 11%



Overfitting validation

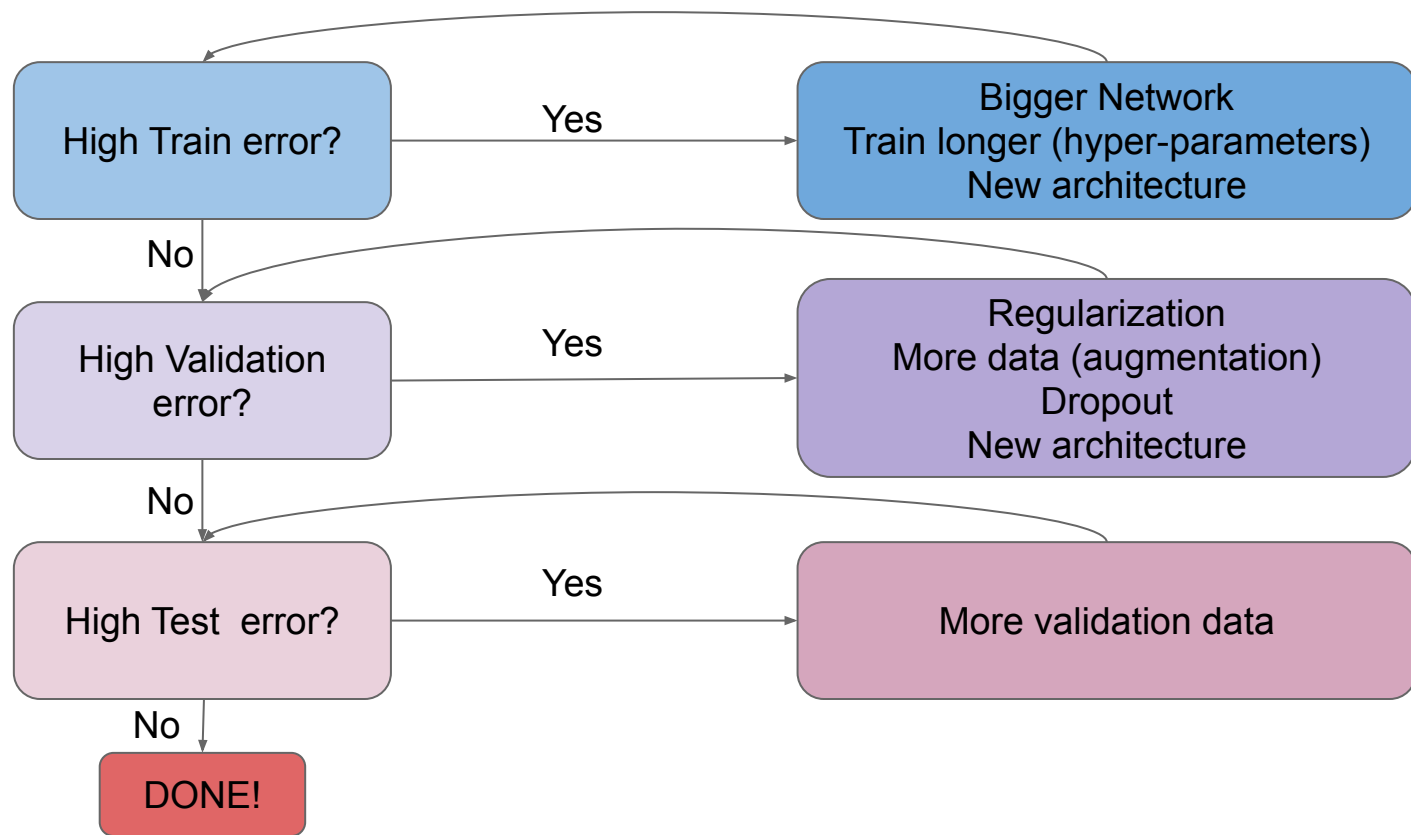
# Strategy for machine learning (5)

TRAINING 60%	VALIDATION 20%	TEST 20%
-----------------	-------------------	-------------


Human level error	.	.	1%
Training error	.	.	1.1%
Validation error	.	.	1.2%
Test error	.	.	1.2%



# Strategy for machine learning (5)



# References



Nuts and Bolts of Applying Deep Learning by Andrew Ng

<https://www.youtube.com/watch?v=F1ka6a13S9I>



# Questions?



**UNIVERSITAT POLITÈCNICA DE CATALUNYA**  
**BARCELONATECH**

**Department of Signal Theory  
and Communications**

*Image Processing Group*